

AIRFLOW PIPELINES, ENGINEERED

Airflow: DAG power, none of the daemon.

Orchestration that scales with you.

+117k

downloads

<50 MB

RAM, total

SQLite

WAL state

THE SHIFT

Where it hurts today

Three things quietly cost teams time and trust every single week.

A metadata DB to babysit

DAGs drift from reality

CI for DAGs is painful



wpipe turns these into code you control.

What you gain



Code-first

Pipelines are Python — versionable, testable, reviewable.



Resume, don't restart

Every step checkpoints to SQLite WAL.



Featherweight

<50 MB RAM. No broker, no daemon.

Crash at step 7? Resume from step 4.





Every step persists its state to SQLite WAL — a rerun cost you never pay again.

THE ARCHITECTURE

One pipeline, fully instrumented

State, retries, parallelism and audit — all inside the same run.



wpipe vs. the status quo



Code-first

Pipelines are Python — versionable, testable, reviewable.



Resume, don't restart

Every step checkpoints to SQLite WAL.



Featherweight

<50 MB RAM. No broker, no daemon.

Based on what the engineering community hits daily.

Answers four questions in one run ↘

Did it run, how far, what failed, and can it resume? Yes, yes, exactly, and automatically.

```
python my_pipeline.py
```

SEE IT IN ACTION

Six lines. That's the learning curve.

```
from wpipe import Pipeline, step

@step(name="job", retry_count=3, retry_delay=2)
def job(data):
    return {"ok": run(data["job"])}

pipe = Pipeline(pipeline_name="airflow_like")

CHECKPOINTS · RETRIES · TIMEOUT · PARALLEL · TRACKER
```



Code-first

Pipelines are Python — versionable, testable, reviewable.



Resume, don't restart

Every step checkpoints to SQLite WAL.



Featherweight

<50 MB RAM. No broker, no daemon.

Try wpipe in your next pipeline

Install the library, wrap your step, and watch checkpoints and docs appear by themselves.

```
pip install wpipe
```